

Penalized Cox Regression

Simulated data

```
library(survival)
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 3.6.1
```

```
## Loading required package: Matrix
```

```
## Loading required package: foreach
```

```
## Loaded glmnet 2.0-18
```

Data generation

```
set.seed(1234)

N <- 1000
p <- 30 # total num. features
nzc <- p/3 # num. 'true' predictors, one third of them B are not 0

X <- matrix(rnorm(N * p), nrow = N, ncol = p)
beta <- rnorm(nzc)
y <- local({
  linear_predictor <- X[, seq_len(nzc)] %*% beta
  hazard <- exp(linear_predictor)
  y_time_event <- rexp(N, rate = hazard)
  y_time_censoring <- rexp(N, rate = hazard * 0.5)
  y_time <- pmin(y_time_event, y_time_censoring)
  y_event <- y_time_event <= y_time_censoring
  Surv(y_time, y_event)
})
```

We should have about 1/3 of the data points censored:

```
table(y[,2])
```


```
##
##    0    1
## 346 654
```

Fit the model

$$\alpha \sum |\beta_i| + (1 - \alpha) \sum \beta_i^2 \leq c$$

how do we choose λ ?


```
fit <- glmnet(X, y, family = "cox")  
plot(fit)
```



Selecting a threshold through cross-validation

the lower the better, 25 coefficient are not 0, but it's better take a more conservative one (15,16)

```
set.seed(1234)  
fit.cv10 <- cv.glmnet(X, y, family = "cox")  
plot(fit.cv10)
```



```
#we can also see the numeric value
str(fit.cv10)
```

```
## List of 10
## $ lambda      : num [1:52] 0.361 0.329 0.3 0.273 0.249 ...
## $ cvm         : num [1:52] 13.7 13.6 13.5 13.5 13.4 ...
## $ cvsd        : num [1:52] 0.0537 0.0515 0.0528 0.0538 0.0546 ...
## $ cvup        : num [1:52] 13.7 13.6 13.6 13.5 13.5 ...
## $ cvlo        : num [1:52] 13.6 13.5 13.5 13.4 13.4 ...
## $ nzero       : Named int [1:52] 0 1 1 1 2 3 4 5 5 5 ...
##   .. attr(*, "names")= chr [1:52] "s0" "s1" "s2" "s3" ...
## $ name         : Named chr "Partial Likelihood Deviance"
##   .. attr(*, "names")= chr "deviance"
## $ glmnet.fit:List of 12
##   ..$ a0        : NULL
##   ..$ beta       : Formal class 'dgCMatrix' [package "Matrix"] with 6 slots
##     .. .. ..@ i      : int [1:757] 0 0 0 0 3 0 3 7 0 3 ...
##     .. .. ..@ p      : int [1:53] 0 0 1 2 3 5 8 12 17 22 ...
##     .. .. ..@ Dim     : int [1:2] 30 52
##     .. .. ..@ Dimnames:List of 2
##       .. .. .. ..$ : chr [1:30] "V1" "V2" "V3" "V4" ...
##       .. .. .. ..$ : chr [1:52] "s0" "s1" "s2" "s3" ...
##     .. .. ..@ x      : num [1:757] 0.065 0.1252 0.181 0.243 0.0447 ...
##     .. .. ..@ factors : list()
##   ..$ df        : int [1:52] 0 1 1 1 2 3 4 5 5 5 ...
##   ..$ dim       : int [1:2] 30 52
```

```
## ..$ lambda : num [1:52] 0.361 0.329 0.3 0.273 0.249 ...
## ..$ dev.ratio: num [1:52] 0 0.00581 0.0107 0.01484 0.02202 ...
## ..$ nulldev : num 7698
## ..$ npasses : int 1845
## ..$ jerr : int 0
## ..$ offset : logi FALSE
## ..$ call : language glmnet(x = X, y = y, family = "cox")
## ..$ nobs : int 1000
## ..- attr(*, "class")= chr [1:2] "coxnet" "glmnet"
## $ lambda.min: num 0.00725
## $ lambda.1se: num 0.0243
## - attr(*, "class")= chr "cv.glmnet"
```

Estimated coefficients:

```
# we can extract one of the threshold
coef(fit.cv10, s = "lambda.1se") #lambda.min global minimum, lambda.1se more conservative one
```

```
## 30 x 1 sparse Matrix of class "dgCMatrix"
##      1
## V1  2.013506525
## V2  0.009425558
## V3  0.358727691
## V4  1.286244321
## V5  0.488818732
## V6  0.948241038
## V7 -0.202498677
## V8 -1.059839099
## V9 -0.221302386
## V10 -0.961357414
## V11 .
## V12 .
## V13 .
## V14 0.038965359
## V15 .
## V16 .
## V17 .
## V18 .
## V19 0.012206575
## V20 .
## V21 .
## V22 -0.045944015
## V23 .
## V24 .
## V25 .
## V26 .
## V27 .
## V28 .
## V29 0.019593248
## V30 0.002423260
```

Compare estimated values with ‘true’ values:

```

beta.true <- c(round(beta, 2), rep(0, ncol(X) - length(beta)))
beta.est <- round(coef(fit.cv10, s = "lambda.1se"), 2)
cbind(beta.true, beta.est)

```

```

## 30 x 2 sparse Matrix of class "dgCMatrix"
##      beta.true      1
## V1      2.49  2.01
## V2      0.05  0.01
## V3      0.46  0.36
## V4      1.58  1.29
## V5      0.62  0.49
## V6      1.19  0.95
## V7     -0.28 -0.20
## V8     -1.35 -1.06
## V9     -0.29 -0.22
## V10    -1.18 -0.96
## V11      .    .
## V12      .    .
## V13      .    .
## V14      .    0.04
## V15      .    .
## V16      .    .
## V17      .    .
## V18      .    .
## V19      .    0.01
## V20      .    .
## V21      .    .
## V22      .   -0.05
## V23      .    .
## V24      .    .
## V25      .    .
## V26      .    .
## V27      .    .
## V28      .    .
## V29      .    0.02
## V30      .    0.00

```

Make predictions

Using built-in functions:

```

predict(fit.cv10, newx = X[1:5, ], s = "lambda.1se")

```

```

##      1
## [1,] -5.1652836
## [2,]  0.6997955
## [3,]  2.9578906
## [4,] -1.4213800
## [5,]  2.0290144

```

‘Manually’:

```

b <- coef(fit.cv10, s = "lambda.1se")
b.i <- which(b != 0)
bnz <- b[b.i]
y0 <- X[1:5, b.i, drop = FALSE] %*% bnz
print(y0)

```

```

##           [,1]
## [1,] -5.1652836
## [2,]  0.6997955
## [3,]  2.9578906
## [4,] -1.4213800
## [5,]  2.0290144

```

Case study nr. 1

Load data and packages

```

library(survival)
library(glmnet)

load('./LymphomaData.rda')

str(patient.data)

```

```

## List of 3
## $ x      : num [1:7399, 1:240] -0.221 -0.1786 -0.0503 -0.1922 -0.2944 ...
## $ time   : num [1:240] 5 5.9 6.6 13.1 1.6 1.3 1.4 2.2 3.4 2.2 ...
## $ status: num [1:240] 0 0 0 0 1 1 1 1 1 1 ...

```

```

gex <- t(patient.data$x)
y <- Surv(patient.data$time, patient.data$status)


```

Exploratory analysis


```

plot(survfit(y ~ 1))


```



```
hist(as.vector(gex))
```




```
hist(colMeans(gex))
```

```
hist(apply(gex, 2, sd))
```

Histogram of apply(gex, 2, sd)



Split the data randomly into a training and a testing set

```
i.training <- sample.int(nrow(gex), size = 160, replace = FALSE)
i.testing <- setdiff(seq_len(nrow(gex)), i.training)


gex.training <- gex[i.training,, drop = FALSE]
y.training <- y[i.training,, drop = FALSE]

gex.testing <- gex[i.testing,, drop = FALSE]
y.testing <- y[i.testing,, drop = FALSE]
```

Train the model

```
fit.cv10 <- cv.glmnet(gex.training, y.training, family = "cox")

plot(fit.cv10)
```



```
b <- coef(fit.cv10, s = "lambda.min")
sum(b != 0)
```

```
## [1] 9
```

```
round(b[b != 0], digits = 3)
```


```
## <sparse>[ <logic> ] : .M.sub.i.logical() maybe inefficient
```

```
## [1] 0.114 0.399 0.008 -0.078 -0.039 -0.016 0.008 -0.014 0.078
```

Test the model

```
score.testing <- predict(fit.cv10, newx = gex.testing, s = "lambda.min")
hist(score.testing)
```

Histogram of score.testing



Question: how well is the score predicting survival?

A continuous predictor vs a right-censored time-to-failure outcome: Cox regression!

```
summary(coxph(y.testing ~ score.testing))
```

```
## Call:
## coxph(formula = y.testing ~ score.testing)
##
##   n= 80, number of events= 47
##
##               coef exp(coef) se(coef)      z Pr(>|z|)
## score.testing  2.5045   12.2369   0.7892  3.173  0.00151 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##               exp(coef) exp(-coef) lower .95 upper .95
## score.testing    12.24    0.08172    2.605    57.47
##
## Concordance= 0.624  (se = 0.046 )
## Likelihood ratio test= 9.6  on 1 df,  p=0.002
## Wald test               = 10.07  on 1 df,  p=0.002
## Score (logrank) test = 10.29  on 1 df,  p=0.001
```

```
#higer score higher risk,
```

A more interpretable scale: by IQR variation:

```
score_scaled.testing <- score.testing / IQR(score.testing)
summary(coxph(y.testing ~ score_scaled.testing))


## Call:
## coxph(formula = y.testing ~ score_scaled.testing)
##
##      n= 80, number of events= 47
##
##              coef exp(coef) se(coef)      z Pr(>|z|)
## score_scaled.testing 0.6758    1.9656   0.2130 3.173 0.00151 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##              exp(coef) exp(-coef) lower .95 upper .95
## score_scaled.testing    1.966    0.5087    1.295    2.984
##
## Concordance= 0.624  (se = 0.046 )
## Likelihood ratio test= 9.6  on 1 df,   p=0.002
## Wald test            = 10.07  on 1 df,   p=0.002
## Score (logrank) test = 10.29  on 1 df,   p=0.001
```

We can split the scores into 2 categories, and compare patients with 'low' vs 'high' score:

```
gex_risk <- ifelse(score.testing <= median(score.testing), "low", "high")
table(gex_risk)
```

```
## gex_risk
## high low
##    40  40
```

```
fit.KM <- survfit(y.testing ~ gex_risk, conf.type = "log-log")
plot(fit.KM)
```



```
fit.KM
```

```
## Call: survfit(formula = y.testing ~ gex_risk, conf.type = "log-log")
##
##               n events median 0.95LCL 0.95UCL
## gex_risk=high 40      28   3.8    2.1    6.6
## gex_risk=low  40      19  11.6    4.9    NA
```

are they significantly different? the p value is low so they are significantly different

```
survdif(y.testing~gex_risk)
```

```
## Call:
## survdiff(formula = y.testing ~ gex_risk)
##
##               N Observed Expected (O-E)^2/E (O-E)^2/V
## gex_risk=high 40      28    18.7    4.58    8.05
## gex_risk=low  40      19    28.3    3.04    8.05
##
## Chisq= 8 on 1 degrees of freedom, p= 0.005
```

Question: how is the 6 months survival for patients classified as low risk, compared to patients classified as high risk?

```
summary(fit.KM, time = 6)
```

```
## Call: survfit(formula = y.testing ~ gex_risk, conf.type = "log-log")
##
##           gex_risk=high
##      time      n.risk  n.event   survival   std.err
##      6.000      14.000   24.000     0.381     0.079
## lower 95% CI upper 95% CI
##      0.229      0.531
##
##           gex_risk=low
##      time      n.risk  n.event   survival   std.err
##      6.0000     22.0000  13.0000     0.6682     0.0755
## lower 95% CI upper 95% CI
##      0.4975     0.7923
```

AUC is high it's good

```
library(survivalROC)
ROC <- survivalROC(Stime = y.testing[, 1],
                  status = y.testing[, 2],
                  marker = score.testing,
                  cut.values = quantile(score.testing, prob = 0:10/10),
                  predict.time = 10,
                  method = "KM")
ROC$AUC
```

```
## [1] 0.7307709
```